

# Seminar Cloud Computing

## From Concept to Production: Enablements TinyML in Industrial Setting

Trung Nguyen  
Technische Universität München

November 2024

### Abstract

In recent years, Artificial Intelligence (AI) and Machine Learning (ML) have received tremendous amount of attention in both industry and research world. However, conventional Machine Learning demands high computing capability which limits its usage to only larger computing units. The paradigm shift to Tiny Machine Learning (TinyML) is revolutionizing industries by enabling the deployment of machine learning models on low-power, resource-constrained devices. Being one of the most rapid developing field of Machine Learning, TinyML promises to benefits multiple industries. However, building a production-ready tinyML system poses different unique challenges. In this paper, we explore the key obstacles faced when developing and deploying TinyML models in production environments, including model optimization, hardware limitations, software integration, and maintaining performance in real-world conditions. Additionally, we present real-world use cases of TinyML in industrial settings, showcasing its transformative impact. We also discuss practical approaches and strategies presented by recent researches [20] to overcome these challenges, providing insights into how TinyML systems can be successfully scaled and implemented in production.

## 1 Introduction

Traditional Machine Learning Models, especially Deep Learning Models typically require substantial amount of computing capability to operate effectively. These models are often trained on powerful Graphics Processing Units (GPUs) and produce large models ranging from tens or hundreds of gigabytes (GB) down to smaller models in the range of 10 to 100 megabytes (MB). However, the memory

requirements during runtime for these models far exceed what microcontrollers (MCUs) can handle. The paradigm shift to TinyML is driven by the prevailing number of Microcontroller Units (MCU) currently circulating in the industry. According to a recent report [2, 1], as of 2021, around 31 billion MCUs were shipped worldwide annually. The MCU market size is projected to increase in upcoming years [2]. This creates a big incentive for researchers and industry players to put effort into developing the technology further. TinyML aims to enable data processing or inferencing directly on embedded systems, particularly on Internet Of Things (IOT), instead of streaming to the cloud. TinyML models can operate on energy- and memory-constrained devices by minimizing communication with external servers, using optimized architecture designs, and employing compression techniques such as quantization and pruning. The advancement of tinyML has positively influenced multiple industries and sectors such as: industrial IOT [18], healthcare [8], agriculture [24], IOT in smart-city [11, 18].

Although various organizations have been actively investing resources into machine learning and data science projects, only around 13% of these projects successfully progress to production.<sup>1</sup> This low success rate highlights the complex challenges and gaps in translating ML concepts into fully functional, production-grade systems, especially in the realm of TinyML. In this paper, we aim to address these challenges and outline what is necessary to transition a TinyML project from concept to production, ultimately creating value for industrial applications. In Chapter 2, we provide an in-depth overview of the fundamental concepts, techniques, and development

---

<sup>1</sup><https://venturebeat.com/ai/why-do-87-of-data-science-projects-never-make-it-into-production/>



Figure 1: TinyML Deployment pipeline

pipeline of TinyML. Following this, in Chapter 3, we draw on recent research to investigate the specific challenges encountered in developing and deploying TinyML models in production environments, as well as practical strategies to overcome these obstacles. Next, in Chapter 4, we explore real-world applications of TinyML in industrial settings, illustrating its transformative potential. Finally, in Chapter 5, we consider the future trajectory of TinyML, discussing upcoming advancements and their implications for broader industrial adoption.

## 2 TinyML Overview

### 2.1 Key Concepts and Techniques

Pruning removes unnecessary neurons or connections from a neural network to reduce its size and complexity. This optimization lowers memory and computational demands, making the model faster and more efficient for deployment on resource-limited devices like microcontrollers.

Quantization reduces the precision of model parameters (e.g., from 32-bit to 8-bit), which decreases model size and computational needs. This technique enables efficient model inference on devices with limited memory and processing power, such as those used in TinyML applications.

With the TinyML model running on edge, the data is stored, processed and analysed internally rather than at an external server or cloud.

### 2.2 TinyML pipeline

**Data Collection:** The workflow starts with collecting data from an IoT device. This data serves as the input for training machine learning models. The IoT device can gather various types of data

depending on the application, such as sensor data in smart homes or environmental monitoring.

**Preprocessing:** After collecting data, it moves to the preprocessing stage. Here, the data is cleaned and prepared for model training. Preprocessing can involve tasks such as normalization, handling missing values, or feature extraction to improve the quality of the input data.

**Model Training:** The preprocessed data is fed into an ML framework, where a machine learning model is trained. This stage involves using machine learning algorithms to find patterns in the data, building a predictive or analytical model that can generalize from the data.

**Model Conversion:** Once the model is trained, it is converted into a format suitable for deployment on resource-constrained devices, like microcontrollers. This step might involve techniques such as quantization (reducing the precision of model weights) or pruning (removing unnecessary parts of the model) to reduce the model’s size and computation needs.

**Model Packaging:** After conversion, the model is packaged into a deployable form. This includes bundling the model with any necessary runtime components to allow it to be executed efficiently on the target device.

**Deployment:** The next step involves deploying the packaged model onto an IoT device. This means transferring the model to the device, setting it up for real-time inference or operation in the field.

**Inference:** Finally, the deployed model performs inference on the IoT device. Inference refers to the process of using the model to make predictions or decisions based on new data collected by the IoT device. The model operates locally on the device without needing continuous cloud connectivity, enabling real-time decision-making at the edge.

## 3 Enablements of TinyML in Industrial Setting

### 3.1 Challenges

The fundamental pipeline for deploying a TinyML model to a microcontroller, as described above in Chapter 2, provides a general framework for model development and deployment. However, in industrial settings, this pipeline often falls short due to a variety of challenges unique to production environments. Firstly, industrial applications typically demand more robust and adaptable systems, requiring continuous model updates and seamless integration within complex IoT infrastructures. Secondly, these demands introduce specific obstacles, such as the inability to regularly update models on resource-constrained devices and the difficulty of integrating TinyML systems with existing IoT architectures.

#### 3.1.1 Adaptation to unseen scenarios and constantly changing conditions

In traditional machine learning (ML), models can be regularly updated with new data through a constant internet connection, allowing them to stay relevant as conditions evolve. This capability, however, is often unavailable for TinyML models deployed on microcontroller units (MCUs) in IoT or edge devices, which typically lack continuous internet access. Consequently, once a model is deployed on an MCU, it may remain static for extended periods, operating without the benefit of regular updates.

This static deployment significantly impacts TinyML models over time, especially in dynamic environments where data patterns evolve. In real-world applications, conditions rarely remain constant, and the input data the model receives can change due to factors such as seasonal variations, environmental shifts, or user behavior changes. Without regular updates, a static model can struggle to adapt to these evolving patterns, a challenge known as concept drift (when the relationship between input and output changes) and data drift (when the distribution of input data changes). Both forms of drift can lead to a decline in model performance, making it essential to consider ways to address these challenges in TinyML deployments.

#### 3.1.2 Integrating with Existing IOT System and Management of TinyML Models in Heterogeneous Devices

Integrating TinyML-enabled microcontrollers (MCUs) into existing IoT systems presents several unique challenges. IoT infrastructures are often built on established architectures, protocols, and data formats that may not be directly compatible with TinyML MCUs. Adapting these MCUs to work within the existing IoT ecosystem requires significant customization to ensure seamless data exchange, communication, and interoperability.

Moreover, existing IoT systems frequently rely on centralized processing or cloud-based analytics, while TinyML emphasizes local computation on edge devices. This architectural shift introduces complexities in synchronizing data and managing hybrid workflows that combine cloud and edge processing. Industrial IoT deployments also come with strict security and compliance requirements, so integrating TinyML MCUs necessitates additional security protocols, device authentication, and data encryption to maintain data integrity across the IoT network.

Maintenance and scalability issues can complicate the integration of TinyML into IoT systems at scale. Updating models on distributed MCUs, managing firmware compatibility, and monitoring device health require robust infrastructure and automated processes, adding layers of complexity to deploying and managing TinyML solutions within existing IoT frameworks. Moreover, to manage and operate such sophisticated system requires deep technical capability which many of professionals with good domain lacks. This urges not only the need for effective methods to manage and maintain TinyML systems at scale but also the importance of making these systems accessible to a broader range of professionals. Enabling more individuals to work with TinyML systems will be crucial for ensuring widespread adoption and operational success in industrial environments. [11, 16, 9, 22, 23].

### 3.2 Bringing tinyML to production environment

In order to target each of the outlined challenges, we have found in recent researches the following approaches proposed by Ren et al. [20, 21, 19].

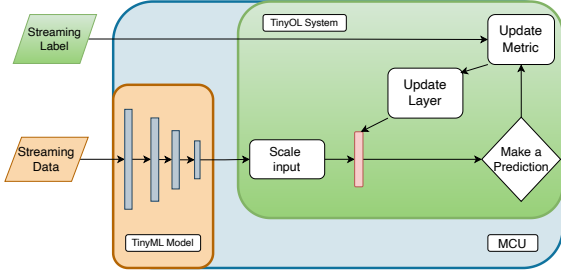


Figure 2: TinyOL Components on MCUs (adapted from Ren et al.)

### 3.2.1 Advanced On Device Learning Methods

Creating a general NN model for general use is impractical for MCUs due to their resource-constraint nature. A simple NN model can easily exceed the hard storage room of a MCU. Moreover, for a general NN model to function, it needs to be updated with field data on the fly, which usually not possible as MCUs do not possess constant connection to the internets, especially while operating at edge.

In their paper on TinyOL (Tiny Machine Learning with Online Learning), Ren et al. introduced a novel system that enables on-device learning. This approach allows models deployed on embedded devices, such as sensors or microcontrollers, to learn and adapt in real time without requiring external re-training.

The core of TinyOL lies in the additional layer marked in red in Figure 2 referred to as the **update layer** which consists of neurons that can be customized, initialied, and updated on the fly. This layer serves as an output layer that can adapt to new data while preserving the original network’s structure. Following is the breakdown of the other components in Figure 2.

- **Streaming Data:** This component represents the live data input collected by the MCU from its environment. It is fed into the TinyML model continuously, allowing the system to make predictions based on real-time information.
- **TinyML Model:** The core machine learning model deployed on the MCU. It takes streaming data as input and performs initial processing. This model consists of several layers (depicted by vertical bars), which extract features from the data.
- **Scale Input:** After the TinyML model pro-

cesses the input data, it passes through a scaling component. This step standardizes or normalizes the data to ensure it fits the model’s expected input distribution, helping maintain consistent performance as new data arrives.

- **Make a Prediction:** After scaling and processing, the TinyOL system uses the model to make predictions. This decision node represents the point at which the system outputs a result based on the latest input data.
- **Update Metric:** When true labels (ground truth) are available, the system can update evaluation metrics based on the prediction’s accuracy. This component helps measure model performance and provides feedback for updating the model’s parameters.
- **Streaming Label:** This component represents the true labels or ground truth data available intermittently. When available, it is used to calculate metrics and adjust the update layer, enabling the model to correct itself over time and adapt to changing data patterns.

**TinyOL Workflow** The **TinyOL workflow** is summarized in the authors’ algorithm [20]. The workflow proceeds as follows:

1. *Initialization:* TinyOL initializes both the TinyML model and the TinyOL system for on-device learning.
2. *Streaming Data Processing:* As new data points are received from the data stream, TinyOL processes each sample iteratively.
3. *Normalization and Standardization:* The running mean and variance of the incoming data are updated for standardization purposes, preparing the data for the model.
4. *Prediction and Model Update:* TinyOL uses the standardized input to make a prediction. If a true label is available, TinyOL computes the error between the prediction and the label, updating evaluation metrics. The model’s weights are then adjusted accordingly using techniques such as **stochastic gradient descent (SGD)**.

**Minium Resource Consumption:** After processing each data point, TinyOL discards it to conserve memory. Compared to traditional batch/offline training methods, this efficient approach enables

TinyOL to handle high volumes of streaming data, which is essential for resource-constrained MCUs with limited storage and processing capacity. The framework operates with as little as 7 KB of RAM and 135 KB of Flash memory.

**Continuous Adaptation:** TinyOL’s architecture allows it to continuously adapt to new data patterns over time, addressing issues like **concept drift** and **data drift**. By updating model weights in real time, the system ensures that the deployed model remains relevant and can handle changes in input data characteristics that may not have been present during initial training.

### 3.2.2 Management of TinyML systems at scale

There has been increased attention to TinyML Management with different methods proposed such as: TinyML as a service, Model Card to describe and organize the information of trained models [12], or TinyML Operations (TinyMLOps) [6] to build an abstraction layer for various hardware compilers. However, these approaches reveal certain drawbacks as they lack features for users to store and exchange information about existing resources.

The paper introduces a TinyML low-code management framework called SeLoC-ML (Semantic Low-Code ML), designed to simplify the deployment, management, and scaling of TinyML applications. The essentials of SeLoC-ML can be broken down into two main parts: to build a KG ontology on the existing overall TinyML system and to inject the KG into a low-code platform<sup>2</sup>.

**Construction of the Knowledge Graph (KG):** By reusing existing standards, like the **Semantic Sensor Network (SSN)** and **W3C Thing Description (TD)**<sup>3</sup> ontologies (a structured schema), an ontology that includes definitions and metadata for TinyML models, such as neural network architecture, layer configurations, and hardware requirements can be built. These data then can be stored using a graph database to facilitate management, discovery, and deployment of TinyML applications.

<sup>2</sup>A low-code platform is a development environment that enables application creation through visual interfaces and minimal coding, making it accessible to users with limited programming knowledge. <https://www.ibm.com/topics/low-code/>

<sup>3</sup><https://www.w3.org/WoT/>

Figure 3 is a simplified visual example of a knowledge graph. A total of 6 ML models are shown on the left side of the figure. One model is especially displayed in details. The "Move" model, as depicted in the knowledge graph, represents a TinyML motion-detection or activity-recognition model with defined input and output requirements, metrics, and compatibility conditions. By including this model in the knowledge graph, the TinyML management system can ensure the Move model is deployed on compatible hardware, facilitating real-time motion analysis or gesture detection in various applications. On the right-side cluster represents various IoT devices or sensor types, each represented by green nodes labeled with device names like "arduino\_nano\_ble.33" and "adafruit\_tflite\_micro." Finally, the link from Smart Sensor to Move Model in the knowledge graph represents a data dependency and compatibility relationship. It indicates that for the Move Model to function effectively, it requires input from a device that fits the description of a Smart Sensor.

**Integrating into a low-code platform:** Having the knowledge graph in form of a graphDB contructed, next step is to integrate with a low-code platform, to enable non-experts to manage and deploy TinyML models across diverse embedded devices. There are several different low-code platforms in the market such as: SAP Build<sup>4</sup> or Siemen Mendix<sup>5</sup>. Low-code software development cuts down on coding by using simple, drag-and-drop visual tools. This makes it quick to build, easy to update, and straightforward to manage business applications without needing much programming experience.

This approach addresses the challenges of **updating models on distributed MCUs** and **managing firmware compatibility** by leveraging Semantic Web technologies and a Knowledge Graph (KG) to centralize and standardize information about TinyML models and IoT devices. The framework’s *ontology* provides a standardized way to describe ML models and IoT devices, facilitating tracking of model versions, configurations, and updates across various devices. By storing information about model updates, dependencies, and compatibility requirements in the KG, the system enables automated discovery of compatible models for each device, allowing administrators to push updates more efficiently. The *querying capabilities* of the Knowledge Graph make it easy to identify which MCUs require updates, simplifying

<sup>4</sup><https://www.sap.com/products/technology-platform/build.html>

<sup>5</sup><https://www.mendix.com>

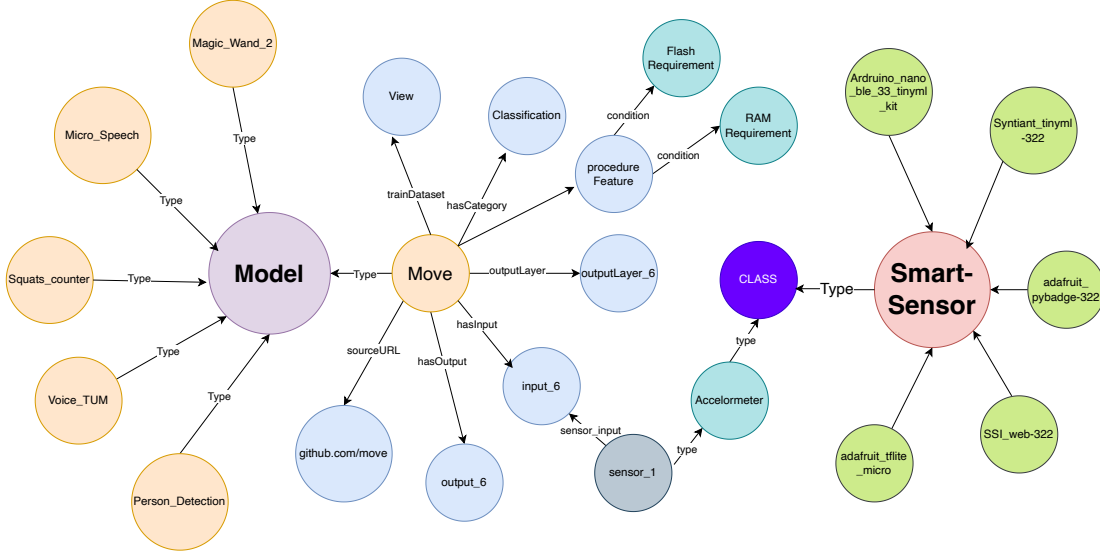


Figure 3: Visualization of System's Knowledge Graph (Simplified)

model deployment across large-scale, distributed systems. Additionally, the KG maintains details about each device's firmware, hardware specifications, and compatibility requirements, ensuring that only compatible firmware and model updates are deployed. This centralized approach minimizes the risk of deploying incompatible firmware or models, which is crucial in environments with diverse MCU configurations.

## 4 Use Cases of TinyML in Industrial Setting

### 4.1 Industrial Predictive Maintenance

Predictive Maintenance (PdM) refers to a practice where business, based on series of historical and failure data, to predict in the future potential health and availability of equipments as well as anticipate defects in advance before they occur, thereby increase equipment lifespan, reduce downtime and operating cost [14]. The practice of constantly observing and predicting has its own inherent flaws and challenges such as financial and organizational constraints, limited source of historical data, and deployment of suitable PdM model [26].

TinyML offers a promising solution to several challenges by enabling intelligent predictive maintenance through embedded, low-power machine learning al-

gorithms that can process data directly on the device. Unlike traditional ML in PdM, which typically relies on centralized data centers and cloud infrastructure, TinyML method helps close the gap in deploying predictive maintenance models by reduce the dependence on large-scale data infrastructure and addressing data source and computational limitations [4]. By allowing edge devices to process and analyze data locally, TinyML reduces the need for extensive data transmission and storage [15]. This enhances the efficiency and reliability of IoT systems by ensuring that only relevant, actionable insights are sent back to the central system.

#### 4.1.1 TinyML vs. Standard ML

According to [5, 26], there are 80% solutions for PdM found in literatures are ML-based and only 20% are TinyML-based. This shows that the TinyML-based solutions in industry are still under-discovered as it still has certain weaknesses compared to Standard ML approaches.

Standard ML models operation requires a lot more computational power, however, as result, are able to handle complex situations at higher accuracy. On the other hand, Standard ML performs with high latency and power consumption, due to the need to transmit data to servers which poses security breaches.

At the edge, where electricity and network connectivity are limited, TinyML proves to be a more viable solution for PdM tasks.

### 4.1.2 Use Cases in PdM

**Anomaly Detection** In extreme environments such as at the ocean where thousands of wind turbines are deployed, TinyML has been applied to detect misbehaviors and early signs of defects in turbines by analyzing turbines' spin behaviors. First, a study shows that novel sensors such as the ISM330IS Inertial Measurement Unit (IMU) from STMicroelectronics enables an onboard Intelligent Sensor Processing Unit (ISPU) that allows some data processing to be computed on-board. In combination with Infrared thermography and continuous line laser thermography techniques to visualize damage in wind turbines, TinyML models can provide real-time analysis and defect detection at the edge without extensive data transmission [17].

Another study showcased a system designed to detect anomalies in submersible pumps at wastewater plants. This system leverages an ESP32DEVKIT microcontroller paired with temperature and vibration sensors. By employing an Isolation Forest model for edge-based anomaly detection, the study highlights the feasibility of retrofitting existing equipment with TinyML capabilities.[7]

In a research that deploys a tinyML model on ESP-WROOM-32 MCU, thermal images of machinery can also be analyzed. Inferred data is then sent via MQTT only upon detection of anomalies to enhance maintenance efficiency. [13]

**Operational Monitoring and Analysis** In extreme environments, continuous monitoring of machinery operations becomes increasingly critical, as maintenance and repair costs tend to escalate over time<sup>6</sup>. TinyML enables continuous monitoring for operational efficiency, safety, or specific conditions indicative of the need for interventions.

A study has been demonstrated for real-time impact localization on thin plastic plates using low-power, resource-constrained IoT devices. This approach leverages data from piezoelectric sensors and employs machine learning models, specifically Random Forest and Shallow Neural Networks, to analyze impact locations. The implementation runs efficiently on an Arduino NANO 33 BLE microcontroller, showcasing the feasibility of deploying such systems on embedded devices. [3].

Another use case involves an ultralow-power Smart

IoT device designed to monitor the activity of handheld power tools in construction environments. Utilizing TinyML for edge processing, the system categorizes tool usage into various operational modes, including transport, no-load, metal drilling, and wood drilling. This categorization enables optimized maintenance, extended tool lifecycles, and enhanced safety measures. The device features temperature, humidity, and acceleration sensors for data collection and employs Bluetooth Low Energy (BLE) and Near Field Communication (NFC) technologies for efficient communication and activation. [10]

### Predictive Maintenance

## 4.2 Smart Agriculture

TinyML is poised to play a crucial role in the future of smart agriculture. By enabling real-time, on-device data processing and decision-making, TinyML is empowering farmers to optimize resource usage, increase productivity, and promote sustainable farming practices. As the technology continues to evolve, we can expect to see even more innovative applications of TinyML in agriculture, contributing to global food security and environmental sustainability. [25] presents a smart agriculture system based on tinyML and why tinyML can solve some of the long-existing challenges of agriculture.

7

### 4.2.1 Precision Agriculture

#### Soil Monitoring

#### Crop Health Monitoring

#### Weather Monitoringg<sup>8</sup>

### 4.2.2 Resource Optimization

#### Smart Irrigation

#### Precision Fertilization

<sup>7</sup><https://github.com/SaiJeevanPuchakayala/FarmWiser>

<sup>8</sup><https://blues.com/use-cases/smart-farming-with-iot-and-machine-learning/>

<sup>6</sup><https://www.theaemt.com/resource/navigating-condition-monitoring-challenges-in-hazardous-areas.html?t>

### 4.2.3 Automated Farming Operations

## 4.3 Healthcare Monitoring

### 4.3.1 Wearable Health Monitoring Devices

**Chronic Disease Management:** TinyML-powered wearables can monitor vital signs such as heart rate, blood pressure, and glucose levels in real-time. This allows for continuous assessment of patients with chronic conditions like diabetes or hypertension, enabling timely interventions and personalized treatment plans.

**Early Detection of Health Anomalies:** Devices equipped with TinyML can detect early signs of health issues by analyzing patterns in physiological data. For instance, changes in heart rate variability or oxygen saturation can be monitored to predict potential cardiac events or respiratory issues.

**Fitness and Activity Tracking:** Fitness trackers using TinyML can provide detailed analysis of physical activities, sleep patterns, and calorie expenditure. This data helps users maintain healthy lifestyles by offering personalized recommendations based on their activity levels. Health care for factory workers - detection whether somebody about to faint

In conclusion, TinyML holds great promise for transforming healthcare monitoring by enabling real-time, personalized health insights through low-power wearable devices. As technology advances, its impact on improving patient care and outcomes is expected to grow significantly.

## 5 Future of TinyML

The future of TinyML holds immense potential as it continues to evolve to meet the demands of diverse industrial applications. Key trends that are expected to drive TinyML’s progress include advancements in model optimization techniques, integration with edge AI, and improvements in hardware capabilities. Enhanced techniques like federated learning for TinyML could allow devices to learn collaboratively without sharing raw data, addressing both data privacy and adaptability in dynamic environments. Additionally, further development in quantization, pruning, and neural architecture search will allow models to be more efficiently deployed on constrained devices, expanding TinyML’s applicability.

The expansion of 5G and the evolution of connectivity standards may also accelerate TinyML adoption, as these advancements could support more reliable intermittent connectivity for edge devices. This would facilitate more frequent model updates and allow TinyML models to synchronize with central systems when needed, even in remote areas. Another future direction is the integration of TinyML with the Internet of Things (IoT) on a larger scale, making real-time analytics accessible across widespread networks of low-power devices. TinyML is also expected to expand in areas like smart cities, environmental monitoring, and autonomous systems, where real-time data processing on edge devices is crucial. As the ecosystem grows, more low-code platforms and management frameworks, such as SeLoC-ML, will likely emerge, empowering a broader range of professionals to develop, deploy, and maintain TinyML systems effectively.

## 6 Conclusion

TinyML represents a paradigm shift in deploying machine learning on resource-constrained devices, offering a feasible solution for real-time, low-power data processing directly on edge devices. This paper has highlighted the critical challenges and considerations in developing and deploying TinyML models in industrial settings, including the limitations posed by hardware, the need for continuous adaptation, and the complexities of integrating with existing IoT systems. Practical approaches such as on-device learning methods, like TinyOL, and management frameworks, such as SeLoC-ML, illustrate the ways in which TinyML can be efficiently scaled and managed to meet industrial demands.

The use cases discussed—spanning predictive maintenance, smart agriculture, and healthcare monitoring—demonstrate TinyML’s transformative potential across industries, providing insights into its real-world applications and impact. While challenges remain, the ongoing advancements in hardware, model optimization techniques, and low-code platforms offer promising solutions to further simplify and enhance the deployment of TinyML systems. TinyML is well-positioned to continue expanding its impact, bringing intelligent edge computing capabilities closer to a wider range of industrial applications and enabling smarter, more efficient systems across the globe.



## References

- [1] Microcontroller market research, 2030.
- [2] Microcontroller Market Size, Share & Trends By 2034.
- [3] Smart Objects | Proceedings of the 20th ACM Conference on Embedded Networked Sensor Systems.
- [4] Youssef Abadade, Anas Temouden, Hatim Bamoumen, Nabil Benamar, Yousra Chtouki, and Abdelhakim Senhaji Hafid. A Comprehensive Survey on TinyML. *IEEE Access*, 11:96892–96922, 2023. Conference Name: IEEE Access.
- [5] Mounia Achouch, Mariya Dimitrova, Khaled Ziane, Sasan Sattarpanah Karganroudi, Rizck Dhouib, Hussein Ibrahim, and Mehdi Adda. On Predictive Maintenance in Industry 4.0: Overview, Models, and Challenges. *Applied Sciences*, 12(16), January 2022.
- [6] Mattia Antonini, Miguel Pincheira, Massimo Vecchio, and Fabio Antonelli. Tiny-MLOps: a framework for orchestrating ML applications at the far edge of IoT systems. pages 1–8, May 2022. ISSN: 2473-4691.
- [7] Mattia Antonini, Miguel Pincheira, Massimo Vecchio, and Fabio Antonelli. A TinyML approach to non-repudiable anomaly detection in extreme industrial environments. In *2022 IEEE International Workshop on Metrology for Industry 4.0 & IoT (MetroInd4.0&IoT)*, pages 397–402, June 2022.
- [8] Mamta Bhamare, Pradnya V. Kulkarni, Rashmi Rane, Sarika Bobde, and Ruhi Patankar. Chapter 14 - TinyML applications and use cases for healthcare. pages 331–353, January 2024.
- [9] Miguel de Prado, Manuele Rusci, Romain Donze, Alessandro Capotondi, Serge Monnerat, Luca Benini And, and Nuria Pazos. Robustifying the Deployment of tinyML Models for Autonomous mini-vehicles, July 2020.
- [10] Marco Giordano, Nicolas Baumann, Michele Crabolu, Raphael Fischer, Giovanni Bellusci, and Michele Magno. Design and Performance Evaluation of an Ultralow-Power Smart IoT Device With Embedded TinyML for Asset Activity Monitoring. *IEEE Transactions on Instrumentation and Measurement*, 71:1–11, 2022. Conference Name: IEEE Transactions on Instrumentation and Measurement.
- [11] Dina Hussein, Dina Ibrahim, and Norah Alajlan. Original Research Article TinyML: Adopting tiny machine learning in smart cities. *Journal of Autonomous Intelligence*, 7:1–14, January 2024.
- [12] Andrew Zaldivar Margaret Mitchell, Simone Wu. Model Cards for Model Reporting. 2019.
- [13] Vítor M. Oliveira and António H. J. Moreira. Edge AI System Using a Thermal Camera for Industrial Anomaly Detection. pages 172–187, 2022.
- [14] Samson Ooko and Simon Karume. Application of Tiny Machine Learning in Predicative Maintenance in Industries. *Journal of Computing Theories and Applications*, 2:131–150, August 2024.
- [15] Samson Otieno Ooko, Marvin Muyonga Ogore, Jimmy Nsenga, and Marco Zennaro. TinyML in Africa: Opportunities and Challenges. pages 1–6, December 2021.
- [16] Aditya Jyoti Paul, Puranjay Mohan, and Stuti Sehgal. Rethinking Generalization in American Sign Language Prediction for Edge Devices with Extremely Low Memory Footprint, February 2021. arXiv:2011.13741.
- [17] Fazle Rabbi. A Novel Approach For Defect Detection Of Wind Turbine Blade Using Virtual Reality And Deep Learning.
- [18] Partha Pratim Ray. A review on TinyML: State-of-the-art and prospects. *Journal of King Saud University - Computer and Information Sciences*, 34(4):1595–1623, April 2022.
- [19] Haoyu Ren, Darko Anicic, Xue Li, and Thomas Runkler. On-device Online Learning and Semantic Management of TinyML Systems. *ACM Transactions on Embedded Computing Systems*, 23(4):1–32, July 2024.
- [20] Haoyu Ren, Darko Anicic, and Thomas Runkler. TinyOL: TinyML with Online-Learning on Microcontrollers. April 2021. arXiv:2103.08295.
- [21] Haoyu Ren, Darko Anicic, and Thomas Runkler. How to Manage Tiny Machine Learning at Scale: An Industrial Perspective. February 2022. arXiv:2202.09113.
- [22] Haoyu Ren, Darko Anicic, and Thomas A. Runkler. The synergy of complex event processing and tiny machine learning in industrial IoT. In

*Proceedings of the 15th ACM International Conference on Distributed and Event-based Systems, DEBS '21*, pages 126–135, New York, NY, USA, June 2021. Association for Computing Machinery.

- [23] A. Navaas Roshan, B. Gokulapriyan, C. Siddarth, and Priyanka Kokil. Adaptive Traffic Control With TinyML. In *2021 Sixth International Conference on Wireless Communications, Signal Processing and Networking (WiSPNET)*, pages 451–455, March 2021.
- [24] Vasileios Tsoukas, Anargyros Gkogkidis, and Athanasios Kakarountas. A TinyML-based System For Smart Agriculture. pages 207–212, November 2022.
- [25] Vasileios Tsoukas, Anargyros Gkogkidis, and Athanasios Kakarountas. A TinyML-based System For Smart Agriculture. pages 207–212, March 2023.
- [26] Tiago Zonta, Cristiano André da Costa, Rodrigo da Rosa Righi, Miromar José de Lima, Eduardo Silveira da Trindade, and Guann Pyng Li. Predictive maintenance in the Industry 4.0: A systematic literature review. *Computers & Industrial Engineering*, 150:106889, December 2020.